

Inside Transducers

- **Speaker:** Rich Hickey
- **Conference:** Clojure/Conj 2014 - Nov 2014
- **Video:** <https://www.youtube.com/watch?v=4KqUvG8HPYo>

[Time 0:00:00]

slide:

[Clojure logo]

Inside Transducers

+ `more.async`

Rich Hickey

Hi. Thanks. Thanks for coming. And thanks for using Clojure. Every year this gets bigger, this event, and this community, and this thing that has become Clojure, and it could not be more exciting to see this, and see all the old friends here. And a lot of new people have been introducing themselves to me and telling me what they are doing, and that is always very exciting. So it is really fantastic to see what people are accomplishing, people who are learning, people who are really killing it with Clojure, succeeding in businesses, helping other businesses or starting their own businesses. It is just totally fantastic.

This year is especially exciting for me because, finally, someone has been able to come to the Conj who has been critical in Clojure's development. I think it is very difficult to understand what it takes to make Clojure. And I am not talking about my own activity, but just what is involved in sort of dropping what you are doing, and spending your retirement money, and having your spouse watch all that and be like, "Go for it!" and be supportive of that, and to listen endlessly to minutiae about Clojure and transducers, and all this stuff that nobody should ever have to listen to.

[Audience laughter]

So I am really happy that my wife Stephanie was able to come today, and I think she deserves a big round of applause – Clojure's biggest supporter.

[Audience applause]

Without whom there would be no Clojure. I can promise you that.

All right! So today we are going to talk a little bit about inside transducers – this is not the same talk I gave at Strange Loop – and also a little bit about some new stuff we are doing in `core.async`.

[Time 0:02:18]

slide title: Transducers

[Two photos, one a close-up of a guitar's strings near the bottom of the fret board, and another of a speaker.]

So transducers are great, because you could probably put a new picture up every time you talk about them, and it would still be a valid transducer. Actually, these are much better examples of transducers than burritos are.

Transducers



Figure 1: 00.02.18 Transducers

[Audience laughter]

[Time 0:02:32]

slide title: What are They?

- + extract the `_essence_` of `map`, `filter` et al
- + away from the functions that transform sequences / collections
- + so they can be used elsewhere
- + recasting them as `_process transformations_`

So I am actually not going to spend a lot of time describing what transducers are, but just as a show of hands, how many people saw the Strange Loop talk? All right, everybody else, you are going to have to go back and fill in some more of the details, because I do not have enough time to both explain the ideas behind transducers and the implementation stuff that I really want to focus on today.

There will be three slides now that are in common with that talk. What is the idea behind a transducer? It is that we were writing “map” and “filter” once again for `core.async` and realizing that probably was not the best thing to be doing. Is there some way to extract this stuff out? Finding an example of how to do that in the reducers work, by basically saying “map” and “filter” are just functions that transform another function, some reducing function. And they could be actually reducing, or it could be something that has the same shape as reducing.

And the way to think about that overall is that “map” and “filter” can be thought of as functions that change some process in a particular way, so you can define them as process transformations. And as soon as you do that, you can use them for push. You can use them for pull. You can use them for data. You can use them for events. I think it is a super important abstraction away from context, which allows you to now express things and use them in different contexts, which is sort of what we keep trying to do.

[Time 0:04:00]

slide title: What Kinds of Processes?

- + ones that can be defined in terms of a `_succession_` of `_steps_`
- + where each step `_ingests_` an `_input_`
- + building a collection is just one instance
- + `_seeded left reduce_` is the generalization

So when we talk about being able to modify a process, what kind of process? I mean, there are a lot of processes with a lot of different shapes, and it ends with the transducers can transduce one particular kind of process, which is pretty generic.

Which is any process that can be defined as a succession of steps, where each step takes some new input and sort of incorporates it in the process. Building a collection is an example of that kind of a process. You have a collection so far, and you have one new thing, and you put it, and you add it to the end. But it is just one example.

And I think too much thinking about “map” and “filter” has been connected to collection processing, or list processing, or sequences, for too long, and that is not the bottom. I think the bottom is the step, not the stuff. So seeded left reduce is the generalization of this.

[Time 0:04:55]

slide title: Why 'transducer'?

- + `_reduce_`
 'lead back'
- + `_ingest_`
 'carry into'
- + `_transduce_`
 'lead across'
- + on the way back / in, will carry `_inputs_` across a series
 of `_transformations_`

You know we like our words in Clojure. So “reduce” means “to lead back”. I am not going to talk much about the ingestion part in the rest of this talk, but “transduce” means “to lead across”, and that is definitely what transducers do. They are leading the input across the series of transformations on the way to, well, maybe reduction or maybe something else, some other process. But they lead them across a transformation process, so the name makes sense.

[Time 0:05:27]

slide title: Transducers

- + `(map f)`
- + `(filter pred)`
- + `(take n)`
- + `(mapcat f)`
- + `(partition-by f)`

`_return_` transducers

In Clojure, moving forward, we are going to have transducers be returned by all the sequence functions that you are used to. So “map” of `f` is no longer – it can still take a collection and do what it always did. But when it is not passed a collection, it returns a transducer. So it returns a function that takes a step-like function, a reducing function, and returns one. So “map” `f` returns a transducer. “filter” with a predicate returns a transducer. That transducer modifies the process. A filtering transducer takes the inputs and maybe does not use them sometimes. That is what filtering is.

Mapping is taking the inputs and applying a function to them before you use them. It is modifying a process. And similarly “take”, and “mapcat”, and whatever. They all return transducers. So “map” is not a transducer. “map” of `f` returns a transducer.

[Time 0:06:27]

slide title: Implementing Transducers

- + take and return a 3-arity fn
- + arity-0 flows through to wrapped fn
- + arity-1 is used for completion, if none flow through
- + arity-2 is reducing step, you can morph input, call nested (or not),
 expand etc

What I want to talk about today is sort of the insides. Maybe hopefully you will leave this talk feeling comfortable being able to implement a transducer, or being able to implement a transducing context.

So what does it take to implement a transducer? We implement them in Clojure as a function that takes a 3-arity function, and returns one. That is the transducing job. The step functions themselves, the function you have to return, and the one you can be expected to be passed, have three arities.

The first one is optional, arity-0. When it is used at the bottom, it is used to fabricate the initial result. So you can think of plus, when you call it with no arguments, returns the identity for plus, which is zero. And multiplication returns one, which is the identity value for multiplication. So if that is present, we want the transducers to preserve it. So far, that is the only use of arity-0 in transducers, so we just flow it through.

Arity-1 is used for completion. And if you do not have any completion to do – a lot of transducers will not. You do not have any accumulated state. You are not keeping track of something as you go along, like every step is its own world – then you have nothing to do with arity-1.

In other words, somebody says, “We have this value. Do you have anything more you need to do with it, because we have no more input? I am calling you with just the return value. Do you have anything to do with it?” And usually your answer will be no, and so you just call through to the function that you are wrapping.

[Time 0:08:07]

And the big one is the arity-2 function, which takes the result so far, and that is a very generic idea. That could be anything. It could be nothing, too. There are reducing functions and step functions that do not use that. They just return nil and you are just passing nil around, and nobody knows better from that.

But one way to think about it is that there is this something that is standing in for the process. It may be the result that is being built up. It may be something the process needs in order to do the next step. Whatever that is, is passed along with a new input. And your job as a transducer is to preserve the result so far. You cannot mess with that. And maybe you do something with the input.

You can morph it. That is what “map” does. You could ignore it. Not use the function given this input. You say I am only supposed to be using even numbers and, when I see an odd number, I just skip. And you skip by just doing nothing. You do not call the nested function, and you just return.

You can do expansion and various other things in transducers. So the arity-2 is the reducing step. It has the same shape as the function you pass to reduce: result so far, and a new input, returns the result next.

[Time 0:09:21]

slide:

Code

So everybody says I do not have enough code in my talks. So now I have a big slide that says Code.

[Audience laughter]

And you can imagine – no.

[Audience laughter]

[Time 0:09:34]

slide:

```
(defn map
  ([f]
    (fn [rf]
      (fn
        ([] (rf))
        ([result] (rf result))
        ([result input]
         (rf result (f input)))))))
```

That is a little bit tiny, isn't it? I had this all worked out so the page downs take me to the next thing. If I make this bigger, that is not going to work. Can people see this, or does it need to be bigger?

[Audience response]

It is OK. Great. This is an excerpt of "map". This is not all of "map", because it is missing the meat of "map", but it is the transducer returning arity of "map", mostly. But it serves as a nice example.

"map" takes a function. That is the function you are going to apply to every input, so you are going to "map" increment. Increment would be f. That is the ordinary argument to "map", and this is the arity that does not take the collection, so just "map" "inc". That would be f.

Then what "map" is supposed to return then is a transducer. So a transducer is a function that takes a reducing function, and returns one. So "rf" is the reducing function that you are wrapping, in a sense. And you have to return one. Like I said, a reducing function, or the full shape of a reducing function, is an arity-3 function. You must have arity-0, 1, and 2.

[Time 0:10:51]

[Points at arity-0 definition on slide]

"map" does not accumulate anything from step-to-step. Every step of "map" is its own world. I take the input. I call the function on it, and I use that input later. I have no state in between calls so, A, I have nothing to construct. This is really for the bottom, so that this arity is always just call the nested function.

[Points at arity-1 definition on slide]

When I am passed a result to complete, I have nothing to do. "map" has no more work to do, so you just call the reducing function on the result. Maybe they have something to do. It is not up to you.

[Points at arity-2 definition on slide]

And then finally, there is the actual step. And this is the essence of mapping is in this step. Given a result and an input so far, and some function you want to transform, mapping is taking that function, the result so far, and passing a transformed input to that function. Calling f. We imagined we were calling "map" with "inc". "inc" would flow down here. You increment the input before you use the wrapped reducing function.

So we have modified a process. "rf" was a process. We do not understand what it is, but we just changed it a little bit. We said, "if you were getting a number, you are now getting a little bit bigger number", or whatever f does. This is the simplest example I can give of a transducer.

[Time 0:12:09]

slide:

```
(defn filter
  ([pred]
   (fn [rf]
     (fn
       ([] (rf))
       ([result] (rf result))
       ([result input]
        (if (pred input)
            (rf result input)
            result)))))))
```

If we move on – ooh, look at that. It is still working. So the first example is “map”, one-to-one. “filter” is one of the sort of the three cases. “map” is one-to-one. That is sort of the easiest. “filter” needs to use the result, or not. Again we see arity-0. So “filter” takes a predicate. It says, if this predicate is true of the input, do the work; otherwise, do not.

And it is interesting, because that means of describing what filtering does is a process-oriented way of talking about it. It is not like if you are given something, and we are supposed to be filtering, and you are not supposed to use it, then give me nothing. That is a data orientation. That is not really what filtering is. Filtering is ignoring, and that is exactly what we are going to do.

Again, the first two arities are the same as “map”. When we are given a result and an input, we say, “if the input satisfies the predicate, use it; call the nested reducing function with that input”. In other words, it is OK input. It passes the predicate. We use it.

Otherwise we do not call “rf”. Don’t do anything. This is the way any kind of reductive or conditional transducer would work. If you do not need to use the input, you do not call the nested function. You just do not do it. That is the pretty much the most efficient way to do nothing, is by doing nothing.

[Audience laughter]

[Time 0:13:36]

slide:

```
(defn take
  ([n]
   (fn [rf]
     ;; create state anew each time called
     (let [nv (volatile! n)]
       (fn
         ([] (rf))
         ([result] (rf result))
         ;; note we can't lazily take 0 - needs input to run logic
         ([result input]
          (let [n @nv
                nn (vswap! nv dec)
                result (if (pos? n)
                          (rf result input)
                          result))]
            (if (not (pos? nn))
                ;; stop early, but don't double-wrap
                (ensure-reduced result)
                result))))))))))
```

So that is what we do.

I am not going to dig totally into “take”, but I wanted to show you an example of a stateful transducer. So “take” has some logic that crosses across execution. As you go, you are only supposed to use so many of the inputs, and then you are not supposed to use anymore. But I wanted to show you a few things just so you have some rules of thumb about creating a stateful transducer.

I think there have been some blog posts that sort of made a lot of stateful transducers. The fact that we can do stateful transducers is great. That does not make stateful transducers sort of the most important thing about transducers, or where the emphasis should be.

When you have a stateful transducer, if you have some local state, make sure you create it anew every time you are asked to modify the reducing function. So you will note that this “let” that creates a volatile variable here is inside, is underneath, the call. So every time you are asking to transform a new reducing function, there is a new state associated with that. That is so that independent transduction stacks are independent, including the fact that their state is independent.

[Time 0:15:03]

Like the others, there is no accumulation in “take”. “take” does keep track of a counter, but it is not accumulating values that need to be flushed out during completion. So both init and complete are just dummy implementations.

But there are some interesting things. In particular, you can only be so lazy with a transducer. So “take” 0 cannot be lazy, because none of the logic of a transducer runs until an input has been supplied. So there has to be some input supplied.

But here it is the basic logic of just saying, “Are we done yet?” The one other thing you want to see here is that: “take” is something that aborts the process, so we are going to take until a certain point. Then we are not going to go anymore. The way we do that is with reduced, which is something we already had in Clojure as a way of saying, “stop reduction”. We are saying transduce and transduction and transducers are sort of an abstraction of reduction. But we still use reduced as a way to signal that we want early termination of this process.

And ensure-reduced is a nice way to say: we called our f here and we might have gotten a reduced value back. We do not want to “reduced” it more. We just want to make sure it has just got a single wrapper. So ensure-reduced is another helper that is present I wanted to call your attention to.

All right, so that is “take”.

[Time 0:16:51]

slide title:

```
(def mapcat
  ([f] (comp (map f) cat)))

;; note cat *is* a transducer
(defn cat
  "A transducer which concatenates the contents of each input, which
  must be a collection, into the reduction."
  [rf]
  ;; helper that surfaces reduced from reduce
  (let [rrf (preserving-reduced rf)]
    (fn
      ([] (rf))
      ([result] (rf result))
```



```
([result input]
  (reduce rrf result input))))))
```

“mapcat” is the other kind of transducer. So we have seen one-to-one. We have seen elision, maybe we will not use the input. We have seen early termination in “take”. And now we have expansion.

What does “mapcat” do? It says: I have a function. I am going to call it on something, and presume that that result is itself a collection, and I am going to integrate that collection in the result.

Now “mapcat” itself sort of streams the answer into a resulting sequence, but “mapcat” the transducer sort of concatenates into the transduction, which may be sequence building or something else – whatever is produced by calling f on an input. So you call f on an input; you get a collection. What you want to say is: all of that collection now gets incorporated as input. That is what we want it to feel like.

So we get to see an example of transducer composition here. “mapcat” is just a composition of “map” and “cat”. If we “mapcat” f, we get “map” f composed with “cat”. “cat” is a new thing. We have not seen “cat” before because – I do not know if “cat” is really an idea with sequences alone, but it certainly is one with transducers.

So “cat” is a little bit different in that the “cat” function is a transducer. It is not a function of some argument that returns a transducer. “cat” *is* a transducer, because it has nothing else to do.

So it just takes a reducing function, and takes its input and reduces it into the result. So it is going to take a reducing function, and take input, which is presumed to be a collection, and incorporates it as input. Which means it uses each piece of that collection as an argument to the reducing function. We are just going to reduce, with the reducing function, all those inputs.

[Time 0:18:59]

We got one collection of five things as an input. We took each of those five things and used them as an input to the function we are wrapping. That is what catenation is.

And so “cat” is a useful transducer. Although, once you have “mapcat”, you will probably end up using that alone, but you can use “cat” directly. If you do not have a function that returns a collection – you just have collections – you can just “cat” them if you are composing transducers, and so “cat” is kind of cool as a transducer.

So this gives you expansion. “cat” gives you the ability to say: I had one input, and it yielded multiple inputs to the next stage. So it expanded.

So we had one-to-one. Filtering is potentially reductive. Taking is abortive, and catenation is expansive. That is sort of the whole flavor of things. We can be one-to-one. We can have fewer things than we got. We can have more things than we got. And we can not use everything that we get.

[Time 0:20:09]

slide:

```
(defn partition-by
  ([f]
    (fn [rf]
      (let [a (java.util.ArrayList.)
            pv (volatile! ::none)]
        (fn
          ([] (rf))
```

```

([result]
(let [result (if (.isEmpty a)
                 result
                 ;; flush
                 (let [v (vec (.toArray a))]
                   ;; clear first!
                   (.clear a)
                   ;; unwrapping helper
                   (unreduced (rf result v)))))]
  ;; complete
  (rf result)))

```

OK. And the final and most complex one I want to show you is partition-by. Again, I do not really want to have you think through and understand all the implementation of this, but I just want to point out a couple things.

Again, it has state, so we are going to create that state anew each time we are asked to wrap a reducing function. This is the first case in which we are going to be building up some intermediate results that need to be flushed during completion. So this is the first transducer I am showing you that has an implementation body of any merit inside complete.

So partition-by, if you do not recall, is a function that takes a function, and it calls it on every input. And as long as it is returning the same value, it keeps that stuff in the same result collection. And as soon as it returns a different value, it yields that collection that it has been building up so far and starts building up a new collection. So it sort of snips your data into collections every time the value of f returns a different result on the input that it is passed.

So eventually, if the process that you are transforming terminates – now it may not. There are kinds of transformations that never stop, in which case partition-by is going to yield as this function changes, but if it never changes again, you will never get that last segment.

But if it is a transduction process that terminates, like reduce, it is given a finite collection, it will terminate. It will end. That means that this arity will get called saying, “OK, there is no more input. Time to wrap it up.”

And partition-by says, “Whoa! I have leftovers. The last time this function returned a value, I snipped it off. I returned a collection. Now I have been building up a new collection, and you are telling me we are done. What am I supposed to do with this stuff?”

[Time 0:22:08]

And the answer is, “now is your chance”. You get one more use of the result. You can use the 2-arity version as much as you need. So it ends up that what is going to happen in partition-by is: it is going to take the collection it has been building up. And you do not have to see the details of that. But it is going to pass it to the wrapped reducing function one more time with a value.

I have one more value. I have been building up these partitions. I made one last partition. I am going to call you with the two-argument version saying, “here is the results so far and here is a value”. So that is going to be typical if you have been building up results. You are going to make one more, or maybe you will reduce, but you are going to make a set of calls to the step function, the two-argument step function, in order to flush your results. And you have to be aware of the fact that that may itself call return reduced saying, “I have seen enough from you. I do not want any more input.”

And then, finally, you have to complete. If you have been building up a result, and you are asked to complete, you can call the wrapped reducing function as many times as you want, ordinarily, to

flush out what you have been building up. And then you must call the reduced argument on that, the reduced flavor or arity of this to complete with the final result. OK? So that will be typical of a transducer that builds up over time.

[Time 0:23:49]

slide:

```
([result input]
(let [pval @pv
      val (f input)]
  (vreset! pv val)
  (if (or (identical? pval ::none)
        (= val pval))
      (do
        (.add a input)
        result)
      (let [v (vec (.toArray a))]
        (.clear a)
        (let [ret (rf result v)]
          ;; when nested done, we are too, don't add
          (when-not (reduced? ret)
            (.add a input))
          ret)))))))))
```

And then some tips on your actual step function. So partition-by is also interesting in that it is incorporating input, and it is building this interim result. It ends up every time it produces an output, it clears the array it has been building up so far. But it may, when it has passed that down, have the function it is wrapping say, “I have seen enough.” It returns reduced.

At that point, we do not want to add any more to the input. In other words, if you are accumulating, as soon as the function under you – the function you are transforming – has told you it is done, it has seen enough – in other words it has terminated early – you should not accumulate anymore. You should put yourself in a state so that when you are asked to complete, you say, “I do not have anything to flush,” because you know the function underneath you does not want to see it. It does not want to see any more from you. It said, “I am done”. So that is what this bit does here. So that would be typical of that.

Again, no one is forcing you to write stateful reducing functions or transducers. They are tricky, so you can just use the ones other people write. But if you are interested in doing it, these are the caveats and tips.

[Time 0:25:11]

slide:

```
(defn transduce
  ([xform f coll] (transduce xform f (f) coll))
  ([xform f init coll]
   (let [xf (xform f) ;; transformation of reducing step
         ret (reduce coll xf init)]
     ;; completion
     (xf ret))))
```

All right. Look, that was code.

[Time 0:25:17]

slide title: Transducible Contexts

- + transduce
- + into
- + sequence
- + channels

OK, so that is writing a transducer. What about writing a transducible context? What does that even mean?

Well, a transducible context is some thing, some process that can incorporate a transducer. We are adding a bunch of these to Clojure. One is a new function called “transduce”. It is analogous to reduce, but it takes a transducer, and it will transduce the reduction process with whatever transducer you pass it.

“into” has been enhanced to take a transducer. “into” used to just say, “dump this collection into that collection”, pour it in. Now you can pour it in and stick a transducer in the middle of that, so everything coming from one collection gets transduced and then put in.

“sequence” will take a collection and produce a lazy view of that collection, having pulled it through a transducer.

And finally, channels will accept transducers.

So these are all examples of transducible contexts. Again, it is even less likely you are going to write your own transducible context, but you should be able to, and hopefully seeing the insides of these will help you.

[Time 0:26:22]

slide title: Implementing a Transducible Context

- + select a step fn - [result input] -> result
- + transform it with transducer
- + call per input
 - might be data source or event
- + iff you have a notion of 'done'
 - + call arity-1

So what do you need to do if you are implementing a transducible context? You have to figure out some way to talk about what you are doing as if there was a step function. Now a lot of times there will be a step function, and it will be plain what it is. And other times, you may have to just think a little bit and say, “Is there a way for me to think about what I am doing as if it was a function of some result so far and a new input, producing a result next?” As long as you can do that, you can transduce it. So select a step function with that shape.

Then you are going to have to have some ability to say, “Give me a transducer”. It is your job as a transducible context to apply the transducer to your step function. That should never be a user space activity. Why is that?

[Audience response]

What is a transduced function like? It has got this whole stack of stuff. All right, what happens if you transduce with partition-by?

[Audience response]

[Time 0:27:28]

That transduced step function is now potentially a stateful thing, which means you do not want more than one person to have that, which means you should never hand it to somebody else. You should never hand a transduced reducing function around, which means all transducible contexts should accept a transducer and, inside them, you call the transducer on the step function and make the transformed step function. And you just keep that to yourself.

Then, every time you have a new input, maybe you are proactive. You have got a collection in hand. Your job is to reduce it, so you just have your stuff already.

Or maybe you are reactive. A channel is reacting to stuff coming in from somewhere else. Well, every time there is a new thing, you are just going to call the step function, the transform step function, on it. Maybe it is data source or it may be an event.

If, and only if – you do not have to do this – if you have a notion of being done. If you are processing a collection, you will have the notion of being done. If you are dealing with events from the outside world, you may not have a concept of being done. You are like an infinite process. You are done when they pull the plug on the computer or something.

But if you do have the notion of being done, when you are done, call the arity-1 to allow any flushing to occur. Call the completing arity.

[Time 0:28:50]

slide title: transduce

- + basic transformation of reducing fn
- + reduce, but adds completion call at end
- + implies reducing fn must support arity-1
 - + `_completing_` - helper that adds it

All right, so let us look inside a couple of these. “transduce” is the most basic. It is just like reduce, except it has a transducer and it transforms the reducing function. It also has to add the completion step at the end. It does imply that the reducing function must support arity-1, which it may not do.

[Time 0:29:10]

slide:

Code

Let us look at code. Do I have code? Oh, look, code.

[Time 0:29:13]

slide:

```
(defn transduce
  ([xform f coll] (transduce xform f (f) coll))
  ([xform f init coll]
```

```
(let [xf (xform f) ;; transformation of reducing step
      ret (reduce coll xf init)]
  ;; completion
  (xf ret)))
```

All right, code. I talked too much about this before showing it.

OK, so what does “transduce” do? It calls reduce. All right, but before it calls reduce, so it supports the same form of reduce here, but it actually has different semantics of reduce.

Who knows what the semantics of reduce are when you call it with a collection and no initial value?

[Audience response]

No one, right. No one knows. It is a ridiculous, complex rule. It is one of the worst things I ever copied from Common Lisp was definitely the semantics of reduce. It is very complex. If there is nothing, it does one thing. If there is one thing, it does a different thing. If there is more than one thing, it does another thing.

It is much more straightforward to have it be monoidal, and just use f to create the initial value. That is what transduce does, so transduce says, “If you do not supply me any information, f with no arguments better give me an initial value.”

Note though, there is no requirement that you support that. You do not have to have that arity. You do not have to allow people to call “transduce” with no initial value, because sometimes there is just no good made up from nothing initial value. Somebody needs to think about an initial value, or supply some inputs to your process to get a starting value. Not everything can be made from nothing. It is easy to come up with zero from nothing. But it is not easy to come up with a channel from nothing, or other kinds of things, event systems. You do not have to support this.

[Time 0:30:43]

So what will happen is if somebody gives you an f and it does not support this arity, it will not work. That is OK.

But the normal form takes an initial value and a collection, so it is just the same shape as reduce, except it has this new first argument: transform. This is the reason. This different treatment of f is the reason why we did not just slap transform into reduce. That is why we have transduce.

So this is just what I said verbally before. First take the function you are supposed to be using. In the case of transduce, f is the step function. It is the reducing function. It is just like reduce. Somebody gave you a step to use. So you just transform it. You have not told anybody else about this.

“xf” is now a private transducer stack that has been constructed around this step function. Then “transduce” just pawns off the job to reduce to do this work, and passes the initial value and the collection.

It gets a return value, but this is another piece that is different. It must call completion. Reduce does have a notion of being done. This collection is going to become exhausted. And therefore, it should call completion, so one last call.

This does imply, however, that the step function support arity-1, which again not all existing step functions do.

[Time 0:32:13]

slide title: transduce

- + basic transformation of reducing fn
- + reduce, but adds completion call at end
- + implies reducing fn must support arity-1
 - + `_completing_` - helper that adds it

So we have a helping function called “completing”, which will take any function that is just a plain reducing function, so it only has the 2-arity: result, input, returning result. And it will add arity-1 with an implementation that is just identity. Given this result, return the same result. So that is a way to sort of take an existing function that is not ready to have complete called on it, and make it possible to do that.

[Time 0:32:43]

slide title: `IReduceInit`

- + `IReduce` does too much (no init)
- + `IReduceInit` still does (eats reduced)
 - eventually a variant that flows reducing out
 - but would be breaking to change `IReduce`

All right.

[Time 0:32:44]

slide:

```
(defn transduce
  ([xform f coll] (transduce xform f (f) coll))
  ([xform f init coll]
   (let [f (xform f)]
     ret (if (instance? clojure.lang.IReduceInit coll)
            (.reduce ^clojure.lang.IReduceInit coll f init)
            (clojure.core.protocols/coll-reduce coll f init))))
  (f ret))))
```

So let me just show you one more thing here, which is the real implementation of “transduce” is a little bit more involved, but I wanted to just introduce this because this is another piece of novelty in Clojure that people are wondering about: what is this reduce init thing?

How many people know about `IReduce`? Not too many. `IReduce` is an interface that supports the semantics of Common Lisp `reduce`. It has two methods: one that is called `reduce` and it takes no init value, and the other is `reduce` that takes an init value. It is a way for a collection to say, “I know how to reduce myself. Instead of asking me for a seq to walk through my stuff, just let me do it because I have a faster way to run through all my own stuff.”

So internal reduce is sort of the Java interface version of `CollReduce`. It already existed, and other people implemented it, but most of the existing implementation of `reduce` is done by `CollReduce`, which has the same exact semantics.

But it ends up that it would be nice to have an interface that just represented the second part of `IReduce`. Remember, I said Common Lisp, who knows what the no init value thing does in `reduce`? Nobody. And so to force everyone to keep implementing both, where the first one has these tricky semantics about no values and one value, was tough.

So IReduceInit only has the version of reduce that takes an initial value, and it is implemented by collections. It is going to come up later, because we are going to see “educe” use it.

So all that is happening inside “transduce” is, it says: I am trying to take advantage of the fastest possible way to reduce you, which is what reduce does now as well.

[Time 0:34:44]

slide title: IReduceInit

- + IReduce does too much (no init)
- + IReduceInit still does (eats reduced)
 - eventually a variant that flows reducing out
 - but would be breaking to change IReduce

All right. So IReduce did too much. It has the no init flavor, which I really do not like, but I cannot just change it, because people have implemented it and they have code that depends on it. So I cannot change its semantics. So all I did was split IReduce into two halves. IReduceInit takes an init, and IReduce derives from it and adds the no init flavor.

But IReduceInit still does something I wish it did not, which is that it eats reduced. In other words, if there is a reduced returned by the function internally, it eats it. It says, “OK, I should stop”, and it stops, but it returns the final result.

It ends up that a lot of times, if you are using reductions in nested contexts, whether you are processing a tree or you are doing transduction, which is sort of a tree of ... or at least a list of nested reductions, you might want to a reduction in an internal step. You need to know if that bailed out, so that you could bail out.

So I would like a variant that actually returns reduced, but we have not had time to integrate that yet. So you see IReduceInit. It is just an interface that implements the reduce version that takes an init.

[Time 0:36:00]

slide title: sequence (LazyTransformer)

- + (sequence xform coll)
- + builds linked list (thus caches)
- + transformed step is list append
- + tricky bit is that a transducer might not yield any output for a particular input
 - + thus to satisfy pull, must loop

All right. Another transducing context is “sequence”. We always had “sequence”. I do not know if anybody ever used it. It basically was just a way to explicitly say: I want to get a sequence out of this collection.

Now it takes a transformer and it has become substantially more interesting, because the collection might not be lazy, but the return value from sequence *is*, and its use of the transducer is also sort of lazy. And we will see what “sort of” means in a second.

So what is this step of “sequence”? Essentially it is building a linked list. So the transform step is sort of list append. We do not really use list append, and usually list append is really expensive, so it must be doing some fancy modifying version that is efficient of list append that actually can add to

the end of a list. And that is what it does. And that is what “lazy-seq” did as well. “lazy-seq” just knew how to attach things to the end of the last cell.

The tricky bit – if you think about making a lazy list out of a transduced collection – is the fact that that collection might sometimes return nothing. Maybe one of the transducers is “filter”, which means I pulled, because I am going to be eating a sequence by pulling. I am pulling. I say: give me something.

If you just consumed one input from the collection and it filtered it out, what would you hand to the user? You have got nothing yet. You actually have to keep going until you see something. So inside, if you look at the implementation of LazyTransformer, which is tricky, you will see that there has to be a loop there, because you have to keep consuming input until you can produce output.

[Time 0:37:52]

slide title: Lazy vs Pushy

- + Transducers are fundamentally push
- + sequence implements pull via introducing inputs
 - + lazy in input consumption, not output production
- + thus sequence not as lazy as lazy-seq
 - but still usefully lazy in practice

The other thing that is tricky about this is: it sort of changes the notion of what it means to be lazy. And I do not really want to change that. I just want you to think about another notion of lazy. Transducers are fundamentally push. You supply input and then the logic runs. They are not fundamentally pull like lazy sequences are.

So “sequence”, the function, implements pull by going around the other end and saying: you wanted something, so I am going to take another input and feed it to the step function. Now we just said one tricky bit is: you may ask for something and I put it through the step function, and the step function gives me nothing, in which case I have got to keep doing this until I have one thing to return to you. But what else could this have done?

[Audience response]

It could have returned one thing each time I ask. It could return nothing. It also could return what?

[Audience response]

More than one thing. Where is that going to go? I asked for one thing. You called the step function. It gave you, like, six things. What is going to happen?

[Audience response]

Yeah, I am just going to attach them to the linked list, which means that it is not quite this laziness we are used to. We are used to saying, “do not run any little part of the calculation until I ask for one thing, and then just do exactly that much work”, which ends up being a huge amount of overhead.

[Time 0:39:11]

So now we are saying you are going to ask for one thing. What is actually lazy is not your consumption, but the production. In other words, we are lazily going to add only as much to this to produce an output to give you, but each step might produce more than one thing, so it is kind of pushy.

So it is lazy in the input consumption, which means if you have something that expands to an infinite sequence, it is not going to work. If you have something that, in a single step, expands to something that consumes all of your memory, that is also not going to work. In particular, it is not going to work

because, even if your result produces something lazy, it is going to be eagerly consumed. *Nothing* inside any of the transducers uses laziness at all, because we want to be able to have these other semantics.

So that is something to be aware of. A lot of people initially trip up over trying to use something lazy in a step, and are surprised that it got fully realized. But that is the nature of this.

I still think this is a useful granularity for laziness. I do not think people are consuming all their memory in a single step. And otherwise, this gives you the best of both worlds, because you get very high performance and only as much caching as you need per step. And it still is sort of windowed. So I think it would work for most things, but it is something to be aware of.

[Time 0:40:34]

slide title: educe / Eduction

```
+ _educ_ - 'lead out'
+ (educ xform coll) -> Eduction
+ IReduceInit / Iterable / Seqable / Sequential
+ Work happens every use!
  no caching, unlike seqs / sequence
+ Subsequent transductions combine work w/o intermediates
```

OK. I do not think I showed any – yeah, I do not show you the code for that.

All right, the other thing that is new in here is “educ”. This is just basically I saw this word and I am like, I have to figure out how to make a language feature out of this.

[Audience laughter]

That is not actually true. There were many, many different names. There was a thing, and then there was the name, not the other way, but it is. Educ, it is so – this is so cool, right?

Educ means “lead out”. We had reduce, which is to lead back, like towards the thing you are making. And we have transduce, which is to lead across this transformation. We have educ, which is just to sort of lead out.

To where? Like where is it going? It does not say. That is what is really cool. It is actually sort of a word game, too, because it is like, if you have ever seen the definition of reduce, it says it takes R and I, and it returns R. It is like reduce without the R. We do not have the – we do not know what we are going to do with this yet. We do not have the result. We are not there yet. We just take the R off. We still end up with an English word, and we are good.

[Audience laughter]

So when you educ a collection, basically you are taking a transducer. You are going to make the recipe for transducing this collection, but you are not going to do anything right then. *Nothing* happens when you call “educ”, except the recipe is created. And you return something that is called an Eduction, or a leading out of this data.

So when does it actually happen? It happens when you use it. It ends up that Eduction implements a bunch of important interfaces. It implements IReduceInit. So it knows how to reduce itself, which means it knows how to transduce itself, and reduce ... and anything that uses IReduceInit.

[Time 0:42:34]

It is iterable. So you can use it anywhere you need something that is iterable. It is seq-able, so you can call “seq” on it, get a seq, and use it in existing code that is expecting lazy sequences. And it is sequential, which is just a tag thing, because it is a sequence of things.

The most important thing about eductions and “educer” is that every time you ask for a seq and walk through it, or every time you ask for an iterator and use it, and every time you transduce or reduce it, every single time the work is going to be done across the entire collection, or as much of it as you consume with the process. There is no caching like there is for seq. So if you had side effects, you would seem them go off over and over again. Don’t have side effects. This is not really what this stuff is for.

[Audience laughter]

So the work happens with every use. And this is one of those tradeoffs. You probably do not want to hand around an education and have ten copies where the work is really expensive. You should probably pour that into a collection and share the results.

But you may very well want to give away an education if it is pretty inexpensive and you do not know how much of it is going to be consumed by each user. Then it is worth doing.

Also, the nice thing is: you can have an education as long as you are not really ready to use it. The cool thing about educations is that if you subsequently transduce them, then the recipes fold together, and it is as if you put them, as if you composed both sets of transducers together. And there is no intermediate stuff created. And you can do that over and over again.

So you can say: I know about this much of the work to do. I had the source material, so I cannot just give you a transducer. I know what the source was. I had the initial set of transformations. But now I am handing it to another stage that I do not want to know about, and it does not want to know about me. It knows it can transduce that some more, hand it to the next guy who can transduce it some more, and then finally someone can pour it into a collection, reduce it, or something like that. And all those intermediate operations are going to fold together, and there will be *no* intermediate data created.

[Time 0:44:42]

slide:

Code

It is just all work. So we can look at this.

[Time 0:44:45]

slide:

```
(deftype Education [xform coll]
  Iterable
  (iterator [_] (.iterator ^java.util.Collection (sequence xform coll)))

  clojure.lang.Seqable
  (seq [_] (seq (sequence xform coll)))

  clojure.lang.IReduceInit
  (reduce [_ f init] (transduce xform f init coll))

  clojure.lang.Sequential)

(defn education
  "Returns a reducible/iterable/seqable application of"
```

```
the transducer to the items in coll. Note that these applications
will be performed every time reduce/iterator/seq is called."
[xform coll]
(Eduction. xform coll))
```

That is what an educt does. This is all of it right here. So Eduction is just a deftype. It implements Iterable. Now the key thing here is it just uses “sequence” to do that. It implements Seqable. It also used “sequence” to do that. It implements IReduceInit. It used “transduce” to do that.

And Sequential is just a tagging interface. So you see there is nothing inside this.

But other tricky aspects of this is: it is not a collection. Unfortunately, java.util.Collection is this giant interface, and satisfying it does not make sense for this class. But it is kind of unfortunate, because Java built a lot of things saying, for instance, all collection constructors take other collections, so they know how to initialize themselves from another collection. It would have been vastly better had they said, “I take an Iterable”, because that is all they really need to know.

So that is a little bit tricky, but we are enhancing, say, “vec”, so that it will be able to quickly and efficiently with both ... using “reduce” will build up a value quickly given an eduction. And “eduction” just makes one of these things, so nothing to it.

[Time 0:46:12]

slide title: channels + transducers

```
+ Being able to apply same transducers to events / async get key value
  prop vs laziness et al
+ transformed step is buffer's add!
  [buf input] -> buf
+ new buffer semantics due to expansion
  if not full?, one input may add more than one item to buffer
```

All right. And finally, we have channels as the last example of transducible contexts. This is really the whole point. A lot of what I have talked about already are variants with some performance enhancements, possibly, over what you can do with sequences and pull.

But the fact that you can use transducers in those contexts, and in completely push contexts like channels, that is the Holy Grail of this thing. That is what this is all about. The fact that you can make a transducer stack and pass it to reduce the most on-demand, do it now, eager thing you can think of, and also pass it to a channel, which does not even have any stuff to do until later when some asynchronous thing comes by with some input.

These are the opposite parts of programming. They are opposite. But transducers connect the two. You do not need a different transducer. You can use the same transducer. So this is the point of transducers, is the fact that they cross across this.

So if we look at transducing a channel, we need to think about channels as if they were a step function. And it ends up in the bottom of channels is this function that looks like this. It takes a buffer and an input, and returns a buffer. It is just a matter of sort of looking at that function and saying, “that is a reducing function”. “reduce” was not called inside channels because it does not have all this stuff. It is getting one thing at a time. But there is this step, and it is called sequentially. It follows all the rules I talked about before, so it is perfect.

[Time 0:47:43]

So what happens in a channel is: this buffer add function gets transformed by the transducer. The one trick of this is the same one we saw earlier with “sequence”. Buffers have a notion of being full. So when you are one away from being full, are you full?

[Audience response]

No. Somebody is going to ask the buffer, “Are you full?”, it is going to say, “No, I have got room”. It used to be, it was saying, “I am not full”, and expecting one more thing to get added, and be asked again, “Are you full now?”

But what happens in a transduction? Well, you are going to get one more input. You are going to let one pending write occur. That is going to go through a transducer, and what could come out?

[Audience response]

More than one thing. And it ends up that buffers, now the semantics are, if you are not yet full, you are going to accept one or more things. And all the buffers have been modified to do that.

It ends up that that is, again, sort of a nice fit. If you are going to consume all of memory, it will not work. But otherwise, the notion – as soon as you have done that with one step, it is going to report full until it drains all the way down to what its target was.

So you cannot use fixed sized buffers, but on the other hand, I think you can implement the logic of saying: the activity of this buffer is gated by its size. It is just not limited to that size at any one point in time. So we have added this semantic to all of those.

[Time 0:49:20]

slide title: More...

- + parallel transduce
- + supportive implementation macros
- + primitives?
- + multi-arity?
- + kv-transduce?
- + functional stateful transducers?

All right. There are many places we want to take this. Certainly I am hoping to get to the point where we are going to take the parallelism that we built for reducers and sort of bring it into this model, so it takes transducers. It should be very straightforward to do.

You saw a lot of repetitiveness in the implementations. A lot of them dummy out: 0-arg, and 1-arg, variants. If you look at the implementation of reducers, there was a macro that made that go away. I will probably implement that, so other implementers will sort of be able to focus on just what is unique about their transducer.

We have ideas about supporting primitives.

There are open questions about multiple arity. “map”, in fact, does support multiple arity, but not all the other ones do, and it does not really make sense for a lot of the others, but it might for “filter”.

Reducers had a notion of being able to fold maps using a key value function as opposed to a function of key and value. That is very efficient. It would be nice to find a way to talk about that and make it explicit.

[Time 0:50:20]

slide title: more.async

- + better integration with existing code
- + channels as promises
- + facilities for endpoints
- + all works-in-progress

Maybe we can talk about that last bit.

All right. I am running out of time, but I just want to take you quickly through some of the new things that are coming into core.async. They are in these categories: better integration with code that already exists, using channels as promises, making it easier and better to use channels for, like, RPC kind of endpoints.

[Time 0:50:40]

slide title: Channels become deref-able

- + Integrates with existing code
- + deref reads from channel
 - will block until available
- + timeout flavor of deref also supported
 - will read `_or_` timeout (not both)

And just all this stuff is in progress. So if you are following along with the alphas, you are seeing some of this stuff happen.

The first thing is that channels are going to become deref-able. This is very useful, again, for existing code. How many people have code where the contract of the code is, “Give me something deref-able”? And you are kind of independent then of whether or not that thing is actually one of the reference types or a promise. There are lots of things that are deref-able. And people have made new deref-able things.

This would make it so that a channel is a valid thing to be deref-able. That will have blocking semantics, so deref will be blocking. It will be like promises that way. And when you deref, it will read from the channel.

But, because they are blocking, these channels will also support the timeout flavor of deref. Everybody know that deref can take a timeout and evaluate a return if the timeout occurred? Yes?

[Audience response]

No? Yeah? Okay. Well, deref can take a timeout, and it is useful when the thing that you are derefing can block.

And so the cool thing about channels when you deref them with the timeout is that they will either successfully read and return that value, or time out, but not both. So they will not time out and then later have a side effect of reading because of some pending read. The read will be killed. You will either get the time out or the read value, but not both. So I think this helps integrate to existing code and has useful semantics.

[Time 0:52:05]

slide title: Promise channels

- + create with `_promise-chan_`
- + write once, read many
- + buffer of one
 - always write-ready
- + first value written 'wins'
 - subsequent dropped
- + all readers unblocked by first write
 - always readable once written

The other thing that is coming are promise channels. Promise is cool, but channels are cooler, because they support alt and things like that. And of course, you can think of a channel as if it was a promise. But there is some more you would want to see to make sure that using a channel as a promise was not a pile of convention that you had to get right over and over again.

So we are going to have a function called “promise-chan”, which will return a promise channel. It is a write-once read-many channel. So as soon as somebody has written a value to it, that value will be returned over and over and over again, permanently.

So it has a buffer of one already by default, so you cannot get that wrong. One of the things people do not want is for you to have a promise that blocks them when they fill it. That is no fun, so we are not going to do that. So that will already be built in.

The first value that gets written wins. That is the same as promises. Subsequent ones will just get dropped.

And the other thing that is really neat about it, compared to using channels raw, is the fact that all of the readers will unblock together.

Multi-channels, like the ones we have, all their readers are racing for the next value. And that is not what you want in a promise. You want everybody who is awaiting that value to succeed, not just the first one or the lucky one to succeed. So that is a new semantic. It is not actually a different semantic of the channel. It is the semantic of the buffer that is used to create it, which is kind of neat. In other words, we implemented promise channels just by making a new kind of buffer. So that is coming.

[Time 0:53:42]

slide title: Endpoints

- + If you give me a channel to fill with a response, how do I know you won't block me?

`_offer!_`

- + prefer over `put!` at ingest

- + and `_poll!_`

only when you can presume item already there, please don't poll with `poll!`

- + neither blocks / parks

return nil if no room / no item

If you think a little bit about that use case of being a supplier to a promise-chan, you get into the third category I wanted to talk about today, which is just being a consumer of channels in a more RPC-like way, or a more API-like way. I think there are sort of two fundamental ways to use channels.

One is you are building a flow network in your program. So channels are the interface points of subsystem boundaries. And you say: this subsystem has an input channel, and that is where it gets its stuff from, as opposed to calling a function on it. And it uses one or more output channels to distribute its results, as opposed to sort of somebody having called the function and getting the results. That decoupling is import for building systems.

That is why we like channels, and you are going to compose a system out of channels orchestrating ... out of processes, orchestrating the channels to wire them together. But you end up with a stable network of relationships that are used to flow data.

But if you think about something like a promise-chan, or using a channel for RPC, that is not that kind of enduring relationship. You *can* use channels for these things, but the semantics get a little bit trickier. In particular, that thing I just talked about for promise channels exists for RPC-like uses of channels.

So how many people are using channels as a way for an asynchronous function to return a value to its caller?

[Audience response]

Yeah. OK. Look, it is already happening. All right. Now how many people are concerned, when they put something on that channel, that it might block them if the person who made the channel messed up and made an unbuffered channel and has not bothered to come back and read it?

[Audience response]

I know I worry about that. And you do not have a great way to communicate about the semantics except, again, via convention to say, “You better give me a buffered channel or dropping channel, because I do not want you to block me.”

[Time 0:55:38]

So we are going to have two new functions. The most important one is offer!. What offer! does is it says, “If there is room for this in the channel, that is going to get written to the channel. If there is not, it is just going to return nil. Did not do it. No. Did not happen.

As an API author, as an author of somebody who is using RPC with channels or filling promises, you are going to tell the user, “I am going to offer you a value.” Now that is a way to tell them quite clearly: they better have room for it, or they better be willing to have it get dropped, because now *you* cannot get blocked. If you call offer!, it will not block, which means you can use it in “go” blocks. It will not block.

And this is definitely a *much* better semantic for those kinds of scenarios. You can say, “I will offer you values,” and that is what your documentation should say. This function will offer the value to the channel you provide. So that is a lot cleaner.

There is a corresponding logic for poll!. This is a lot less justified except for sort of completeness and balance. Here you are saying, “You better have given me a channel with stuff in it already, because I am going to poll it”, and either there will be stuff there, or it is just going to give me nil.

Now, how many people think polling is good?

[Audience response]

Nobody. Right? That is why it is called poll!. If you put this in a loop, you are polling, and everyone will know because you just said, “poll!” I am polling. So I think this should be pretty rare, but I would not say that there is no good use for it, but there is no good use for polling.

So neither of these block.

[Time 0:57:20]

slide title: pipeline

- + "don't do too much work in your callback function" is back:
 - "... in your channel transducer"
 - + runs under the channel lock
 - + inhibits other channel ops
- + use pipeline instead

I am not going to have enough time to talk about “pipeline”, except to say “pipeline” is awesome.

[Audience laughter]

But one thing – well, first of all, transducers and channels are awesome, but you should be aware of the fact that a transducer in a channel runs in the most protected part of the channel code, which is in the part that transfers values in and out of the buffer, which is the part that runs under the lock of the channel, which means that you will impede other activity on that channel for the duration of your transducer, which means we are back to the old, “do not do too much work in your callback function.” So now it is in your channel transducer, because of these two things.

[Time 0:58:07]

slide title: pipeline

- + Moves items from one channel to another
- + subject to a transformation
- + with user-specified parallelism (first arg)
- + always in-order results

So there is a new set of three functions called “pipeline” that I think encapsulate a whole bunch of what is pretty hard work to get right, involved with pipelining.

So just so everybody is aware, what a pipeline does is it takes a bunch of sequential work. This is usually a stream of inputs and then a job you want to do. It is going to turn those jobs into a set of parallelized work, and then return the results in order out through a serial process.

You can think of “pmap” sort of does this, but “pmap” is sequence to sequence. And so “pipeline” is channel to channel, subject to a transformation using a transducer. And you are also able to control the parallelism. You can say, “Use this many threads or that many threads. Do this many at the same time”. So unlike “pmap”, you have a lot of control. It takes a channel, returns a channel, takes a transducer, but it is parallel in the middle.

There are three flavors. I am going to talk about this.

[Time 0:59:03]

slide title: pipeline

- + 3 flavors
 - `_pipeline_` (computational)
 - `_pipeline-blocking_` (a blocking op per item)
 - `_pipeline-async_` (an async op per item)
- + `pipeline` and `pipeline-blocking` take transducer
 - (`pipeline [n to-ch xf from-ch] ...`)

I am going to go five minutes late.

There are three flavors. There is “pipeline” itself, which is for computational purposes – exactly what the transducer was for. So do not block in this thing. Do not call an async function in this thing. You are just here to calculate. It is going to use the right threads internally to do that.

There is “pipeline-blocking”. This says, “In the middle of the step transformation I want to do, I need to block”. Well, this would be catastrophic inside the transducer in a channel. You are blocking in the middle of a channel, and now the channel is blocked, so pipeline is a better place to do that.

And then there is “pipeline-async”. Maybe you want to take everything that comes through this channel and make a web call. You do not know when it is going to come back, and then proceed out the other side.

How many people have written pipelining themselves manually?

[Audience response]

Yeah. How many people got it right the first time?

[Audience laughter]

Yeah. And how many people find them doing it over and over again, right, because it is a hard thing to sort of encapsulate? So I think channels give us the ability to sort of capture a pretty high level pattern as a reified programming construct, so we have pipeline-async as well.

Pipeline and pipeline-blocking both take a transducer. So it takes N, which is the amount of parallelism, the to-channel – where is this stuff going to go? – the transducer, and the from-channel – where is this stuff coming from? “pipeline-blocking” does the same thing.

[Time 1:00:31]

slide title: pipeline parallelism

- + uses appropriate threads for nature of job
 - compute vs blocking vs async
- + because pipeline is parallel, transducers are applied per-item, not across items
- + so can't pipeline stateful transducers, e.g. partition-by
- + can still be filtering or expansive

Something to be aware of is the fact that because pipeline is parallelized. Well, first of all, it is going to do the right thing. For async stuff, it is going to use go blocks, and for compute stuff, it can use – for blocking stuff, it will use thread. So it will do the right thing.

But because it is parallel, you cannot do any stateful transducers. Each of these things, you took work like this and turned it into that briefly and then turned – which way did it come in? This way, then that, and back out, right? It keeps the order.

But, in the middle, it is parallel, so no stateful transducers. If you look at the doc strings for transducers, for all the things that return transducers, they all say “stateful” when they return stateful. You cannot use them in parallel contexts. You cannot use them in “pipeline”. You are not going to be able to use them effectively with the parallelism versions of ... when we move reducers into a transducer compatible format.

Of course, right? You should think, “Of course not”. That is not a problem. They have utility in one context. They do not in another. All right. That is OK. There should still be round holes even though there are square pegs.

Because of that, your transducer is a “one operation at a time” transducer, no state allowed. But they can still be filtering or expansive. We can take one thing and return more than one thing. We can take one thing and ignore it. What we cannot do is accumulate any information across. No state.

[Time 1:02:04]

slide title: pipeline-async

```
+ pipeline-async takes channel filling fn
  (pipeline-async [n to-ch af from-ch]
    _af_ is [input result-ch] -> nil
+ _af_ should return immediately, having launched async work that fills
  result-ch and closes it
```

All right. “pipeline-async” I really will not have time to go into. It is a little bit different. It does not take a transducer, because it needs a way to get an asynchronous result, and it uses channels to do that. It will soon tell you it is going to offer your results to the channel. But basically it is going to give you your input and a channel to put the result. Then it collapses all that stuff and effectively streams it out the other side.

So between the three of these, that should cover all your pipelining needs.

[Time 1:02:33]

slide:

```
+ core.async and transducers have seen great community uptake and feedback
+ keep it coming
+ more cool stuff on the way for 1.7
+ Thanks!
```

[Clojure logo]

To wrap up, how many people are using transducers?

[Audience response]

How many people are using core.async?

[Audience response]

Yeah, so I think both of these have seen a lot of community uptake early on, especially since transducers have still just been in the alphas. And a lot of great feedback, so please keep it coming. It is really great. There is a lot more stuff coming for 1.7. I think Alex is going to run a thing on feature expressions later, which is probably the next other big thing, so look forward to that and other stuff, and rock on.

Thank you.

[Audience applause]

[Time 1:03:07]